

## Task 3 (25 points) NYC Green Taxi dataset

In task 3, you will analyze the **NYC Green Taxi dataset** and may need to use the **Spark Graph Processing** techniques. The **dataset** is stored in files named with different months of year 2023 under the path `//data-repository-bkt/ECS765/nyc_taxi/green_tripdata/2023/`. You can view the files using the **ccc method bucket ls** command.

The description of each field in order for the **dataset** is shown below (a sample CSV file was given to explore/visualise on your machine):

**lpep\_pickup\_datetime:** The date and time when the meter was engaged.

**lpep\_dropoff\_datetime:** The date and time when the meter was disengaged.

**PULocationID:** Taxi Zone in which the taximeter was engaged. Refer to 'taxi\_zone\_lookup.csv' for details.

**DOLocationID:** Taxi Zone in which the taximeter was engaged. Refer to 'taxi\_zone\_lookup.csv' for details.

**passenger\_count:** The number of passengers in the vehicle

**trip\_distance:** The elapsed trip distance in miles reported by the taximeter

**fare\_amount:** The time-and-distance fare calculated by the meter

**extra:** Miscellaneous extras and surcharges. Currently, this only includes the \$0.50 and \$1 rush hour and overnight charges

**mta\_tax:** \$0.50 MTA tax that is automatically triggered based on the metered rate in use.

**tip\_amount:** Tip amount – This field is automatically populated for credit card tips. Cash tips are not included.

**tolls\_amount:** Total amount of all tolls paid in trip

**ehail\_fee:** Dont care column

**total\_amount:** The total amount charged to passengers. It does not include cash tips

**payment\_type:** A numeric code signifying how the passenger paid for the trip. 1= Credit card 2= Cash 3= No charge 4= Dispute 5= Unknown 6= Voided trip

**trip\_type:** A code indicating whether the trip was a street hail or a dispatch that is automatically assigned based on the metered rate in use but can be altered by the driver

**congestion\_surcharge:** Total amount collected in trip for congestion surcharge

**taxi\_type:** Green Taxi

We will use the two fields, “**PULocationID**” and “**DOLocationID**”, from the **dataset**. The table below shows a few sample records of “**PULocationID**” and “**DOLocationID**” from the **dataset**.

| <b>PULocationID</b> | <b>DOLocationID</b> |
|---------------------|---------------------|
| <b>166</b>          | <b>143</b>          |
| <b>24</b>           | <b>43</b>           |
| <b>223</b>          | <b>179</b>          |
| <b>41</b>           | <b>238</b>          |

There is another file named “**taxi\_zone\_lookup.csv**” stored under the path **//data-repository-bkt/ECS765/nyc\_taxi/**. The “**taxi\_zone\_lookup.csv**” file contains the location info of each LocationID in the **dataset**. This table has the following schema:

**LocationID:** integer (nullable = true)

**Borough:** string (nullable = true)

**Zone:** string (nullable = true)

**service\_zone:** string (nullable = true)

In the table below, we show a few sample records of “**taxi\_zone\_lookup.csv**”.

| <b>LocationID</b> | <b>Borough</b>   | <b>Zone</b>                    | <b>service_zone</b> |
|-------------------|------------------|--------------------------------|---------------------|
| <b>1</b>          | <b>EWR</b>       | <b>Newark Airport</b>          | <b>EWR</b>          |
| <b>2</b>          | <b>Queens</b>    | <b>Jamaica Bay</b>             | <b>Boro Zone</b>    |
| <b>3</b>          | <b>Bronx</b>     | <b>Allerton/Pelham Gardens</b> | <b>Boro Zone</b>    |
| <b>4</b>          | <b>Manhattan</b> | <b>Alphabet City</b>           | <b>Yellow Zone</b>  |

As you can see, the “**PULocationID**” and “**DOLocationID**” fields in the **dataset** are encoded with numbers that you can find their counterpart (**LocationID** field) with location information in “**taxi\_zone\_lookup.csv**”. Please, note the following:

1. The LocationID 264 and 265 are **Unknown** in the Borough field; we see the **Unknown** as one of the borough names;
2. In the Borough field, you can see the same **Borough** has a different **LocationID**, it does not matter because the **LocationID** is a unique key;
3. If you see any obscure description (like NV, NA, N/A, etc) in any field, purely see it as valid names.

Please note that the example tables/pictures provided in this task are based on my results. Your results may differ depending on how you process and display the data. When a screenshot is required, in addition to the results, you must also include your running timestamp in your screenshot. A running timestamp is the recorded time (printed at the beginning of each line in the logs) during the execution of a Spark job; you can capture the running timestamp in the logs. The example tables/pictures provided in each question illustrate the expected format (including the running timestamp and your results) for your screenshot.

- 1) (2 points) Load all monthly CSV files from the NYC Green Taxi dataset into a single dataframe and print the total number of entries in the dataframe. You will need to include a screenshot of your results in your report. For example:

```
2024-10-11 18:15:38.421 INFO scheduler.TaskSchedulerImpl: Killing all running tasks in stage 3: Stage finished
2024-10-11 18:15:38.421 INFO scheduler.DAGScheduler: Job 2 finished: count at NativeMethodAccessorImpl.java:0, took 0.811988 s
The number of entries is: 466523
2024-10-11 18:15:38.476 INFO datasources.InMemoryFileIndex: It took 8 ms to list leaf files for 1 paths.
2024-10-11 18:15:38.546 INFO graph.GraphContext: Touching stage() from shutdown hook
```

**Note that the above figures/values may not represent the actual result.**

- 2) (5 points) Define the **StructType** for the vertexSchema and edgeSchema, and construct the edges dataframe and vertices dataframe. Use the “**taxi\_zone\_lookup.csv**” as the vertex information and the “**PULocationID**” and “**DOLocationID**” fields from the **NYC Green Taxi dataset** as the edge information.

**Then,** Display 5 samples of both the edges and vertices DataFrames, ensuring that **field names are not truncated**. The vertices DataFrame should include the fields: “**id**”, “**Borough**”, “**Zone**”, and “**service\_zone**”, while the edges DataFrame should include “**src**” and “**dst**”. Include a screenshot of the results in your report. For example,

```
2024-10-11 18:30:08,957 INFO scheduler.DAGScheduler: Job 2 finished: showString at Native
2024-10-11 18:30:08,995 INFO codegen.CodeGenerator: Code generated in 22.352903 ms
```

| src | dst |
|-----|-----|
| 82  | 196 |
| 7   | 7   |
| 7   | 7   |
| 166 | 74  |
| 236 | 229 |

only showing top 5 rows

```
2024-10-11 18:30:09,084 INFO storage.BlockManagerInfo: Removed broadcast_5_piece0 on 10.
2024-10-11 18:30:09,085 INFO storage.BlockManagerInfo: Removed broadcast_5_piece0 on q2-
```

```
2024-10-11 18:30:09,342 INFO scheduler.TaskSchedulerImpl: Killing all running tasks in st
2024-10-11 18:30:09,343 INFO scheduler.DAGScheduler: Job 3 finished: showString at Native
2024-10-11 18:30:09,371 INFO codegen.CodeGenerator: Code generated in 14.030698 ms
```

| id | Borough       | Zone                 | service_zone |
|----|---------------|----------------------|--------------|
| 1  | EWB           | Newark Airport       | EWB          |
| 2  | Queens        | Jamaica Bay          | Boro Zone    |
| 3  | Bronx         | Allerton/Pelham G... | Boro Zone    |
| 4  | Manhattan     | Alphabet City        | Yellow Zone  |
| 5  | Staten Island | Arden Heights        | Boro Zone    |

only showing top 5 rows

```
2024-10-11 18:30:09,435 INFO spark.SparkContext: Invoking stop() from shutdown hook
2024-10-11 18:30:09,448 INFO server.AbstractConnector: Stopped Spark@b64fa35[HTTP/1.1, [ht
2024-10-11 18:30:09,452 INFO ui.SparkUI: Stopped Spark web UI at http://a2-721a7e927cd7f8
```

**Note that the above figures/values may not represent the actual result.**

- 3) (4 points) Create a graph using the vertices and edges. Show 10 samples of the graph DataFrame with columns “src”, “edge”, and “dst”. **Do not truncate the names of fields.** You need to give a screenshot of your results in your report. For example,

```
2024-10-11 18:35:09,470 INFO scheduler.DAGScheduler: Job 4 finished: showString at NativeMethodAccessorImpl.java:0, took 0.163509 s
2024-10-11 18:35:09,487 INFO codegen.CodeGenerator: Code generated in 11.91692 ms
```

| src   | edge       | dst  |
|---|------------|--|
| [82, Queens, Elmhurst, Boro Zone]                     | [82, 196]  | [196, Queens, Rego Park, Boro Zone]                          |
| [7, Queens, Astoria, Boro Zone]                       | [7, 7]     | [7, Queens, Astoria, Boro Zone]                              |
| [7, Queens, Astoria, Boro Zone]                       | [7, 7]     | [7, Queens, Astoria, Boro Zone]                              |
| [166, Manhattan, Morningside Heights, Boro Zone]      | [166, 74]  | [74, Manhattan, East Harlem North, Boro Zone]                |
| [236, Manhattan, Upper East Side North, Yellow Zone]  | [236, 229] | [229, Manhattan, Sutton Place/Turtle Bay North, Yellow Zone] |
| [75, Manhattan, East Harlem South, Boro Zone]         | [75, 235]  | [235, Bronx, University Heights/Morris Heights, Boro Zone]   |
| [260, Queens, Woodside, Boro Zone]                    | [260, 160] | [160, Queens, Middle Village, Boro Zone]                     |
| [95, Queens, Forest Hills, Boro Zone]                 | [95, 264]  | [264, Unknown, NV, N/A]                                      |
| [244, Manhattan, Washington Heights South, Boro Zone] | [244, 41]  | [41, Manhattan, Central Harlem, Boro Zone]                   |
| [83, Queens, Elmhurst/Maspeth, Boro Zone]             | [83, 7]    | [7, Queens, Astoria, Boro Zone]                              |

only showing top 10 rows

```
2024-10-11 18:35:09,549 INFO spark.SparkContext: Invoking stop() from shutdown hook
2024-10-11 18:35:09,562 INFO server.AbstractConnector: Stopped Spark@7009362[HTTP/1.1, [http/1.1]] {0.0.0.0:4040}
```

**Note that the above figures/values may not represent the actual result.**

- 4) (5 points) Count connected vertices (the **src** to **dst**) with the same Borough and service\_zone. **Do not truncate the names of fields.** The table columns should include '**id**'(src), '**id**'(dst), '**Borough**', and '**service\_zone**'. You need to show the total number of connected vertices meeting the above condition and give a screenshot of 10 samples of that outcome. For example,

```
2024-10-11 18:43:29,182 INFO scheduler.TaskSchedulerImpl: Killing all running tasks in
2024-10-11 18:43:29,183 INFO scheduler.DAGScheduler: Job 4 finished: count at NativeMet
count: 285732
2024-10-11 18:43:29,293 INFO datasources.FileSourceStrategy: Pruning directories with:
2024-10-11 18:43:29,295 INFO datasources.FileSourceStrategy: Pushed Filters:

2024-10-11 18:43:29,740 INFO codegen.CodeGenerator: Code generated in 13.866982 ms
+-----+
|id|id|Borough|service_zone|
+-----+
|82|196|Queens|Boro Zone|
|7|7|Queens|Boro Zone|
|7|7|Queens|Boro Zone|
|166|74|Manhattan|Boro Zone|
|236|229|Manhattan|Yellow Zone|
|260|160|Queens|Boro Zone|
|244|41|Manhattan|Boro Zone|
|83|7|Queens|Boro Zone|
|223|223|Queens|Boro Zone|
|260|260|Queens|Boro Zone|
+-----+
only showing top 10 rows

2024-10-11 18:43:29,805 INFO spark.SparkContext: Invoking stop() from shutdown hook
```

**Note that the above figures/values may not represent the actual result.**

- 5) (4 points) Find the shortest paths from all other vertices (**Other LocationID**) to vertex 1 (**LocationID=1**). Show 10 samples from your result. The table columns should include '**id\_to\_1**' and '**shortest\_distance**'. **Do not truncate the names of fields.** You need to give a screenshot of your 10 samples in your report. For example,

```
2024-10-22 07:57:51,776 INFO
+-----+
|id_to_1|shortest_distance|
+-----+
|147->1|2|
|19->1|2|
|39->1|2|
|71->1|2|
|180->1|2|
|130->1|1|
|66->1|1|
|138->1|2|
|171->1|2|
|170->1|2|
+-----+
only showing top 10 rows

2024-10-22 07:57:51,838 INFO
```

**Note that the above figures/values may not represent the actual result.**

- 6) (5 points) Perform page ranking on the graph dataframe. You will set the 'resetProbability' to 0.17 and 'tol' to 0.01. And sort vertices (id) by descending according to the value of PageRank. Show the top 5 results. You need to give a screenshot of your results in your report. The table columns should include 'id', 'pagerank'. For example:

```
2024-10-22 08:07:18,806 INFO codegen.1
```

| id pagerank |                    |
|-------------|--------------------|
| 264         | 23.733400016313745 |
| 138         | 7.780849414226171  |
| 132         | 7.195856951107596  |
| 55          | 4.471856828199626  |
| 265         | 4.10645225541102   |

```
only showing top 5 rows
```

```
2024-10-22 08:07:18,864 INFO spark.Sp
```

**Note that the above figures/values may not represent the actual result.**